

PERSONAL STATEMENT

The Hong Kong University of Science and Technology

Jianming Xing

MSc in Artificial Intelligence

A learned robot policy becomes useful only when its data, software, and hardware agree about what the robot sees and what it should do. On the CR5, I found that aligned timestamps, consistent state and action fields, and inference from a meaningful robot state could change physical behavior even when the model itself was unchanged. That experience shifted my focus from isolated prediction to the computer-science problems connecting representation, learning, serving, and physical action under changing conditions.

To gain end-to-end competence across machines and computation, I built my education as a bridge rather than a replacement. In June 2025, I earned a B.Eng. in Mechanical Design, Manufacturing and Automation at Harbin Institute of Technology, Weihai. That training grounded me in mechanism behavior, kinematics, and physical constraints. I am now pursuing a second bachelor’s degree in Computer Science and Technology at Harbin Institute of Technology, Shenzhen, expected in June 2027. This expansion lets me reason both about a machine’s physical capabilities and about how algorithms, data structures, interfaces, and runtime decisions determine its behavior. Together, these disciplines help me follow a learned capability from representation to physical execution.

My research examines how action representation shapes what a vision-language-action model can learn. As a Research Assistant at iLearn, HIT Shenzhen, I am second author of Hermite Action Tokenizer for VLA, an unsubmitted working paper, and I am responsible for its discrete autoregressive tokenization branch. I worked on splitting seven-dimensional action chunks into six-dimensional end-effector trajectories and gripper states, using shared breakpoints to remove redundant endpoints and quantization seams while coding gripper changes separately. I also compared adaptive MLP-based and fixed Hermite segmentation. This design raised a central evaluation question: how can I isolate the quality of an encoding from the behavior of a complete policy? In an offline codec benchmark averaging four LIBERO subsets with 8-by-7 action chunks, Hermite’s Endpoint MSE was approximately 99.7% lower than FAST and 82.6% lower than BEAST. The comparison demonstrates codec fidelity under that protocol; it does not establish online rollout success or physical-robot reliability.

Separate systems work on a Dobot CR5 showed how representation choices travel through an implemented learning loop. I anchored HDF5 records to RGB timestamps and aligned robot states, target actions, and gripper feedback through binary search and linear interpolation. Before downstream use, I checked schemas, timing, motion, and workspace constraints, then converted the episodes into LeRobot v2.0 with consistent seven-dimensional robot-plus-gripper states and actions. I adapted OpenPI through a training configuration, WebSocket policy serving, and a CR5 client that assembled observations and executed ServoP arm targets and Modbus gripper commands. In my tests, non-blocking execution increased apparent command frequency, but a new inference could use an intermediate in-motion state before the preceding action chunk had finished, producing back-and-forth corrections. I therefore used blocking, short-horizon ServoP chunks, accepting lower apparent frequency to preserve state-action consistency. This was an execution decision rather than a change in model score. Although Hermite and CR5 were separate efforts, together they taught me to connect an algorithmic question about representation with engineering questions about clocks, software interfaces, observations, and action execution.

Supporting projects extended this perspective. For a weld-seam HMI, I connected

C++/Qt controls to Python/PyTorch segmentation inference through QProcess and visualized robot kinematics in OpenGL, keeping the desktop interface, inference process, and robot model as intelligible components. During the Xbotics Embodied AI Community Internship, I migrated an ANYmal-C navigation task from Isaac Lab into a MotrixLab NumPy environment. I refactored reset and step flows, corrected coordinate handling, and used rollout and reward/termination regression checks to examine whether the migrated environment retained its intended behavior. Both projects made cross-component change a concrete software-design problem.

My immediate goal is AI/ML or learning-systems R&D for embodied agents, building methods that remain dependable when connected to real data and hardware. I want graduate training to deepen my foundations in machine learning, data systems, and scalable software while testing them against sensing and control constraints. After gaining stronger real-system experience, I may consider doctoral study as an option rather than a predetermined step. I seek curricular depth and applied opportunities to turn this perspective into rigorous, transferable engineering. I want to understand how an action representation interacts with the broader intelligence of a model: what a foundation model contributes, how sequential learning changes behavior, and when a representation advantage survives beyond an offline comparison. Hermite and the separate CR5 loop gave me starting cases at opposite ends of that question. I now need stronger AI foundations that connect them without treating a codec result as complete-policy performance. HKUST's MSc in Artificial Intelligence offers that progression.

If the current offerings, including ARIN5204, remain available in my intake and the optional ARIN6900 Independent Project is approved, ARIN5102 AI Fundamentals: Concepts and Methods would consolidate the principles behind my model choices. ARIN5203 Foundation Models and Generative Artificial Intelligence would situate VLA action prediction within a broader model context, while ARIN5204 Reinforcement Learning would deepen the sequential-learning layer. In ARIN6900, I could vary an action representation within a learned policy, compare its predicted trajectories offline, and then study changes in sequential behavior without conflating the two results. This research path would prepare me to develop embodied-AI methods that can progress from controlled comparisons toward physical-agent R&D.